

## 1. Transaction Concepts

A **transaction** is a logical unit of work that consists of one or more database operations (read/write) executed as a single, indivisible unit.

### The SQL Transaction Process

Transactions are controlled using specific SQL commands. The most common commands used for managing transactions are:

1. **BEGIN TRANSACTION:** Starts a transaction.

Example: `BEGIN TRANSACTION TransferFunds;`

2. **COMMIT:** Commits the transaction, saving all changes made.

Example: `DELETE FROM Student WHERE AGE = 20;`

`COMMIT;`

3. **ROLLBACK:** Rolls back the transaction, undoing all changes made since the BEGIN TRANSACTION.

Example: `DELETE FROM Student WHERE AGE = 20;`

`ROLLBACK;`

4. **SAVEPOINT:** Marks a point within a transaction to which you can later roll back.

Example: `SAVEPOINT SP1;`

`//Savepoint created.`

`DELETE FROM Student WHERE AGE = 20;`

`//deleted`

```
SAVEPOINT SP2;  
//Savepoint created.
```

ROLLBACK TO SAVEPOINT Savepoint\_Name;

RELEASE SAVEPOINT Savepoint\_Name;

**Transactional SQL Statement using Explicit Transaction Control in SQL (Structured Query Language).**

**Example 1:**

```
BEGIN TRANSACTION;  
    UPDATE Accounts SET balance = balance - 500 WHERE acc_no = 'A1'; -- Debit  
    UPDATE Accounts SET balance = balance + 500 WHERE acc_no = 'A2'; -- Credit  
COMMIT;
```

If any operation fails, the entire transaction is rolled back to maintain consistency.

**Example 2(Bank Transfer In Detail):**

```
BEGIN TRANSACTION;  
  
-- Deduct $150 from Account A  
UPDATE Accounts  
SET Balance = Balance - 150  
WHERE AccountID = 'A';  
  
-- Add $150 to Account B  
UPDATE Accounts  
SET Balance = Balance + 150  
WHERE AccountID = 'B';  
  
-- Commit the transaction if both operations succeed  
COMMIT;
```

## 2. Properties of Transactions (ACID)

ACID properties ensure **data integrity, reliability, and correctness** in a **transactional database system**. These properties are **essential for maintaining database consistency**, especially in **concurrent transactions and failure scenarios**.

---

### 1. Atomicity (All or Nothing)

#### Definition:

Atomicity ensures that a **transaction is treated as a single, indivisible unit**. Either **all operations of the transaction are executed**, or **none of them occur**.

#### Example:

Consider a **bank transfer of \$500 from Account A to Account B**:

BEGIN TRANSACTION;

UPDATE Accounts SET Balance = Balance - 500 WHERE AccountID = 'A';

UPDATE Accounts SET Balance = Balance + 500 WHERE AccountID = 'B';

COMMIT;

#### Scenario:

- If the **first UPDATE succeeds** (money is deducted from A) but the **second UPDATE fails** (due to a crash or network failure), the transaction **must be rolled back**.
- Without **atomicity**, Account A would lose money without Account B receiving it.

 **Atomicity ensures that either both operations happen or none happen (Rollback on failure).**

---

## 2. Consistency (Valid State Before & After)

### Definition:

Consistency ensures that **a transaction transforms the database from one valid state to another valid state**, maintaining **all integrity constraints**.

### Example:

Consider a **bank balance rule**: **An account cannot have a negative balance**.

BEGIN TRANSACTION;

UPDATE Accounts SET Balance = Balance - 500 WHERE AccountID = 'A';

UPDATE Accounts SET Balance = Balance + 500 WHERE AccountID = 'B';

IF (SELECT Balance FROM Accounts WHERE AccountID = 'A') < 0

ROLLBACK;

ELSE

COMMIT;

### Scenario:

- If **Account A** has **only \$300**, deducting \$500 would make the balance **negative**, violating consistency.
- The transaction **rolls back** to maintain consistency.

 **Consistency ensures that database rules and constraints are never violated.**

---

## 3. Isolation (Concurrent Transactions Do Not Interfere)

### Definition:

Isolation ensures that **concurrent transactions execute independently without interfering with each other**.

### Example:

Two transactions **T1 (Withdraw \$500 from A)** and **T2 (Check balance of A)** occur simultaneously.

### Without Isolation (Dirty Read Problem):

T1:

```
UPDATE Accounts SET Balance = Balance - 500 WHERE AccountID = 'A';
```

T2:

```
SELECT Balance FROM Accounts WHERE AccountID = 'A'; -- Reads incorrect data
```

- If T2 reads Account A's balance while T1 is still in progress, it may see an **intermediate state** (before commit).
- If T1 fails and rolls back, T2 would have used an **incorrect balance**.

✅ Isolation ensures that T2 waits for T1 to complete (via locking mechanisms like Read Committed, Serializable isolation levels).

---

## 4. Durability (Permanent Changes)

### Definition:

Once a **transaction is committed**, its changes **must be permanently stored**, even if the system crashes.

### Example:

Consider a **ticket booking system**:

```
BEGIN TRANSACTION;
```

```
UPDATE Tickets SET Status = 'Booked' WHERE TicketID = 101;
```

```
COMMIT;
```

### 💡 Scenario:

- If a power failure occurs **after the COMMIT**, the changes **must still persist** when the system restarts.

- Database logs (WAL - Write Ahead Logging) ensure durability by saving committed transactions to non-volatile storage.

✓ Durability ensures that once committed, transactions survive failures.

---

## Summary Table of ACID Properties

Property	Definition	Example Scenario
<b>Atomicity</b>	A transaction is either fully completed or fully rolled back	Bank transfer: Deduct money from A <b>and</b> add to B, or rollback both
<b>Consistency</b>	The database remains valid before and after a transaction	A bank account cannot go negative
<b>Isolation</b>	Concurrent transactions execute independently	One transaction updating a balance doesn't interfere with another reading it
<b>Durability</b>	Committed transactions are permanent	A booked movie ticket remains confirmed even after a system crash

---

## Why Are ACID Properties Important?

- ✓ Prevents data corruption and loss
- ✓ Ensures correctness in multi-user environments
- ✓ Maintains integrity even in failures
- ✓ Essential for banking, e-commerce, and enterprise systems

### 3. Serializability of Transactions

**Serial Schedule:** Transactions execute one after another (no interleaving).

**Serializable Schedule:** A non-serial schedule that produces the same result as some serial schedule.

Types of Serializability:

1. Conflict Serializability:

- Two operations conflict if they:
  - Belong to different transactions.
  - Operate on the same data item.
  - At least one is a writer.

Example:

T1: R(A), W(A)

T2: R(A), W(A)

- Non-serial but conflict-serializable if equivalent to  $T1 \rightarrow T2$  or  $T2 \rightarrow T1$ .

2. View Serializability:

- Less strict than conflict serializability.
- Allows blind writers (writes without reading first).

## 4. Testing for Serializability

Method: Precedence Graph

1. Construct a graph where:

- Nodes = Transactions (T1, T2, ...).
- Edge  $T_i \rightarrow T_j$  if an operation in  $T_i$  conflicts with and precedes an operation in  $T_j$ .

2. Check for cycles:

- If there are no cycles, the schedule is conflict-serializable.

Example:

Schedule S:

T1: R(A), W(A)

T2: R(A), W(A)

T3: R(B), W(B)

Conflicts:

- $T_1 \rightarrow T_2$  (both access A, T1 writes first)
- No conflict between T3 and others.
- Graph:  $T_1 \rightarrow T_2$  (acyclic)  $\Rightarrow$  Conflict-serializable.

### Summary

- ✓ Transactions ensure logical grouping of operations.
- ✓ ACID properties guarantee reliability.
- ✓ Serializability ensures correct concurrent execution.
- ✓ Precedence graph helps test conflict serializability.

Next Lecture: Concurrency Control Techniques (Locking, Timestamping, Optimistic Methods)